

Gaussian Naive Bayes Classifier

May 20, 2024

```
[ ]: import numpy as np
      from scipy.stats import multivariate_normal
      from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
      import matplotlib.pyplot as plt

[ ]: # Data Preprocessing
      TRAIN = './mnist_train.csv'
      TEST = './mnist_test.csv'

      # Function to load data and split into features and labels
      def load_data_and_split(filepath):
          data = np.genfromtxt(filepath, delimiter=',', skip_header=1, dtype='int')
          X = data[:, 1:] # All columns except the first one
          y = data[:, 0] # Only the first column
          return X, y

      # Load and split the training and testing data
      X_train, y_train = load_data_and_split(TRAIN)
      X_test, y_test = load_data_and_split(TEST)

      # Normalize the training and testing data
      X_train = X_train.astype(np.float64)
      X_test = X_test.astype(np.float64)

      X_train /= 255
      X_test /= 255

[ ]: class GaussianNaiveBayes:
      def __init__(self):
          self._mu = dict()
          self._sigma = dict()
          self._pi = dict()

      def fit(self, X, y, smoothing=0):
          self._classes = np.unique(y)
          self.n_classes = len(self._classes)

          cov = np.diag(np.diag(np.cov(X, rowvar=False)))
```

```

max_var = np.max(cov)
# Compute mean, var, and prior for each class
for idx, c in enumerate(self._classes):
    X_c = X[y == c]
    self._mu[idx] = X_c.mean(axis=0)
    self._sigma[idx] = np.diag(np.diag(np.cov(X_c, rowvar=False))) + ␣
    ↪smoothing*max_var
    self._pi[idx] = X_c.shape[0] / X.shape[0]

def predict(self, X):
    # Compute posterior probability for each class
    n_samples, n_features = X.shape
    preds = np.zeros((n_samples, self.n_classes), dtype=np.float64)

    for idx, c in enumerate(self._classes):
        prior = np.log(self._pi[idx])
        class_conditional = self._pdf(idx, X)
        posterior = prior + class_conditional
        preds[:, idx] = posterior
    # Return the class with the highest posterior probability
    return self._classes[preds.argmax(axis=1)]

def _pdf(self, class_idx, X):
    mean = self._mu[class_idx]
    sigma = self._sigma[class_idx]
    pdf_values = multivariate_normal.logpdf(X, mean=mean, cov=sigma, ␣
    ↪allow_singular=True)
    return pdf_values

```

```

[ ]: # Initialize Gaussian Naive Bayes classifier
gnb = GaussianNaiveBayes()

smoothing = [0.001, 0.01, 0.1, 1, 10, 100]
accuracy = []
y_preds = dict()

for idx, s in enumerate(smoothing):

    # Fit the model on the training data
    gnb.fit(X_train, y_train, smoothing=s)

    # Predict on test data
    y_pred = gnb.predict(X_test)
    y_preds[idx] = y_pred

```

```

# Calculate accuracy
acc = np.mean(y_pred == y_test)
accuracy.append(acc)
print(s)
print(f"Accuracy: {acc * 100:.3f}%")
print(f"One-Zero Error: {(1-acc)*100 :.3f}%")

y_pred = y_preds[accuracy.index(max(accuracy))]

# Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
plt.show()

```

```

0.001
Accuracy: 78.880%
One-Zero Error: 21.120%
0.01
Accuracy: 78.880%
One-Zero Error: 21.120%
0.1
Accuracy: 78.380%
One-Zero Error: 21.620%
1
Accuracy: 79.800%
One-Zero Error: 20.200%
10
Accuracy: 80.800%
One-Zero Error: 19.200%
100
Accuracy: 79.140%
One-Zero Error: 20.860%

```

